

# Distributed Consensus

*CS 675: Distributed Systems (Spring 2020)*

Lecture 5

Yue Cheng

Some material taken/derived from:

- Princeton COS-418 materials created by Michael Freedman and Wyatt Lloyd.
- MIT 6.824 by Robert Morris, Frans Kaashoek, and Nickolai Zeldovich.
- Utah CS6450 by Ryan Stutsman.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

# Today's outline

1. View changes in primary-backup replication

2. Consensus

- Paxos
- Raft

# Review: Time & Clocks, PB

# With multiple replicas, don't need to wait for all...

- Viewstamped Replication:
  - State Machine Replication for any number of replicas
  - Replica group: Group of  $2f + 1$  replicas
    - Protocol can tolerate  $f$  replica crashes
- Assumptions
  1. Handles crash failures only: Replicas fail only by completely stopping
  2. Unreliable network: Messages might be lost, duplicated, delayed, or delivered out-of-order


# Replica state

1. Configuration: identities of all  $2f+1$  replicas
2. In-memory log with clients' requests in assigned order

$\langle \text{op1}, \text{args1} \rangle$	$\langle \text{op2}, \text{args2} \rangle$	$\langle \text{op3}, \text{args3} \rangle$	$\langle \text{op4}, \text{args4} \rangle$
--	--	--	--


# Normal operation

( $f = 1$ )



# Normal operation


( $f = 1$ )



1. Primary adds request to end of its log

# Normal operation

( $f = 1$ )




1. Primary adds request to end of its log



# Normal operation


( $f = 1$ )



1. Primary adds request to end of its log
2. Replicas add requests to their logs in primary's log order

# Normal operation


( $f = 1$ )



1. Primary adds request to end of its log
2. Replicas add requests to their logs in primary's log order

# Normal operation


( $f = 1$ )



1. Primary adds request to end of its log
2. Replicas add requests to their logs in primary's log order
3. Primary **waits for  $f$  PrepareOKs** → request is **committed**

# Normal operation: Key points

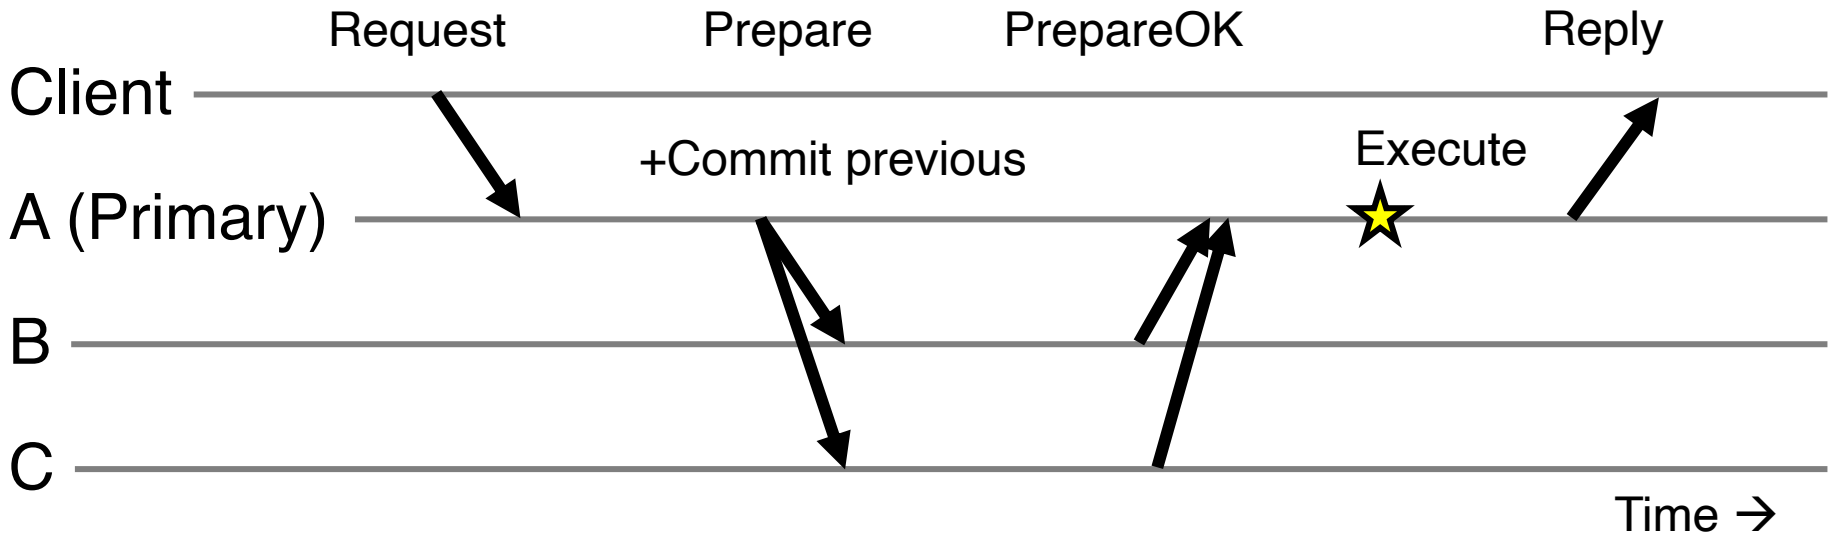
( $f = 1$ )



- Protocol provides state machine replication
- On execute, primary knows request in  $f + 1 = 2$  nodes' logs
  - Even if  $f = 1$  then **crash**,  $\geq 1$  **retains request in log**

# Piggybacked commits

( $f = 1$ )




- Previous Request's commit **piggybacked** on current Prepare
- No client Request after a timeout period?
  - Primary sends Commit message to all backups

# The need for a view change

- So far: **Works** for  $f$  failed backup replicas


# The need for a view change

- So far: **Works** for  $f$  failed backup replicas
- But what if the  $f$  failures include a **failed primary**?
  - All clients' requests go to the failed primary
  - System **halts** despite **merely  $f$  failures**



# Views

- Let **different replicas** assume role of primary over time
- System moves through a sequence of views
  - **View** = (view number, primary id, backup id, ...)





# Correctly changing views

- View changes happen locally at each replica
- Old primary executes requests in the old view, new primary executes requests in the new view
- Want to ensure state machine replication

# Correctly changing views

- View changes happen locally at each replica
- Old primary executes requests in the old view, new primary executes requests in the new view
- Want to ensure state machine replication
- So correctness condition: Executed requests
  1. Survive in the new view
  2. Retain the same order in the new view

# Correctly changing views

- View changes happen locally at each replica
- Old primary executes requests in the old view,

How do they **agree on the new primary**?

What if both backup nodes attempt to become the new primary simultaneously?

2. Retain the same order in the new view

# Today's outline

1. View changes in primary-backup replication

## 2. Consensus

- Paxos
- Raft

# Consensus

- Definition:
  1. A **general agreement** about something
  2. An idea or opinion that is **shared by all** the people in a group

# Consensus used in systems

Group of servers attempting:

# Consensus used in systems

Group of servers attempting:

- Make sure all servers in group receive the same updates in the same order as each other

# Consensus used in systems

Group of servers attempting:

- Make sure all servers in group receive the same updates in the same order as each other
- Maintain own lists (views) on who is a current member of the group, and update lists when somebody leaves/fails



# Consensus used in systems

Group of servers attempting:

- Make sure all servers in group receive the same updates in the same order as each other
- Maintain own lists (views) on who is a current member of the group, and update lists when somebody leaves/fails
- Elect a leader in group, and inform everybody

# Consensus used in systems

Group of servers attempting:

- Make sure all servers in group receive the same updates in the same order as each other
- Maintain own lists (views) on who is a current member of the group, and update lists when somebody leaves/fails
- Elect a leader in group, and inform everybody
- Ensure mutually exclusive (one process at a time only) access to a critical resource like a file

# Consensus

Given a set of processors, each with an initial value:

- **Termination:** All non-faulty processes eventually decide on a value
- **Agreement:** All processes that decide do so on the same value
- **Validity:** Value decided must have proposed by some process



# Safety vs. Liveness properties

- Safety (bad things never happen)
- Liveness (good things eventually happen)



# Paxos

- Safety (bad things never happen)
  - Only a single value is chosen
  - Only chosen values are learned by processes
  - Only a proposed value can be chosen
- Liveness (good things eventually happen)




# Paxos

- Safety (bad things never happen)
  - Only a single value is chosen  **agreement**
  - Only chosen values are learned by processes
  - Only a proposed value can be chosen  **validity**
- Liveness (good things eventually happen)

# Paxos

- Safety (bad things never happen)
  - Only a single value is chosen  **agreement**
  - Only chosen values are learned by processes
  - Only a proposed value can be chosen  **validity**
- Liveness (good things eventually happen)
  - Some proposed value eventually chosen if fewer than half of processes fail
  - If value is chosen, a process eventually learns it

# Paxos

- Safety (bad things never happen)
  - Only a single value is chosen  **agreement**
  - Only chosen values are learned by processes
  - Only a proposed value can be chosen  **validity**
- Liveness (good things eventually happen)
  - Some proposed value eventually chosen if fewer than half of processes fail  **termination**
  - If value is chosen, a process eventually learns it