

Distributed Consensus

CS 675: Distributed Systems (Spring 2020)

Midterm review

Yue Cheng

Some material taken/derived from:

- Princeton COS-418 materials created by Michael Freedman and Wyatt Lloyd.
- MIT 6.824 by Robert Morris, Frans Kaashoek, and Nickolai Zeldovich.
- Utah CS6450 by Ryan Stutsman.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

Midterm

- Wednesday, March 4th, 4:30 – 6:30pm
 - 120 minutes
 - Open-book, open-notes (you may use class notes, papers, and lab materials; you may read them on your laptop, but you are not allowed to use any network)
- Covering topics from lec-1 to lec-5 (but not including Raft)
 - Go-specific debugging questions
 - High-level design questions

Concurrency in Go

- Labs that were completed
 - Possible race condition bugs in Go
 - Go channels
 - Go mutex locks

MapReduce

- Why MapReduce
 - Google workload characteristics
- How MapReduce works
 - Paper
- How data flows within a MapReduce job
 - Use of local file system and use of GFS
- Main shortcomings of using MapReduce

Time & Clocks

- Cristian's algorithm
- Logical Clock algorithm
 - Guarantees if $a \rightarrow b$, then $C(a) < C(b)$
 - How to guarantee a total order of events
 - NYC+SF bank applications of LC – its assumptions: Why it works; under what assumptions it may not work
- Vector Clock algorithm
 - If $V(a) < V(b)$, then $a \rightarrow b$
 - If $V(a) \not< V(b)$ and $V(b) \not< V(a)$, then $a \parallel b$
 - Can use to infer when an event b was aware of / influenced by a

Paxos

- The safety and liveness property of Paxos
- Basic Paxos algorithm
- Basic Paxos examples
- Only proposers knows which value has been chosen
 - If other peers want to know, they must execute Paxos with their own proposal